

## **Assignment 3: Programming Report**

*Lucas Machado*

Student ID: 400507829

Scientific Programming: Operations on Matrices

Monday, March 25, 2024

# 1 Solving the Linear System of equations $Ax = b$

To solve the system of linear equations, we will use Gaussian Elimination. For our algorithm to solve the system, we will use forward elimination and backward substitution. However, we will not use row swapping as it is not necessary to implement, and if we were to implement it, then the code would be more complex, and we would have to deal with more factors. To simplify the algorithm's explanation, we will walk through an example, explain how it works, and provide examples to go with the explanations to make it as easy as possible for anyone to understand. To keep things very simple, we will work with an arbitrary  $2 \times 2$  matrix  $A$  and an arbitrary  $2 \times 1$  vector  $B$  to explain the first steps in the algorithm thus,

Given a matrix  $A$  ( $2 \times 2$ ) and a vector  $B$  ( $2 \times 1$ ):

$$A = \begin{pmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \end{pmatrix}, \quad B = \begin{pmatrix} b_1 \\ b_2 \end{pmatrix}$$

## 1.1 Initial Conditions on Matrices

1. Check to make sure that the matrix  $A$  is square ( $N1 = M1$ ).
2. Verify that the number of rows in Matrix  $A$  is equal to the number of rows in Matrix  $B$  ( $N1 = N2$ ).
3. Ensure that Matrix  $B$  is a column vector ( $M2 = 1$ ).

## 1.2 Creating Augmented Matrix

Create an augmented matrix: by using matrix  $A$  and appending matrix  $B$  to the end of it.

$$\text{Augmented Matrix} = \left( \begin{array}{cc|c} a_{11} & a_{12} & b_1 \\ a_{21} & a_{22} & b_2 \end{array} \right)$$

## 1.3 Normalization of the Pivot Row

We want to normalize each pivot row before we start the forward elimination. This means we will divide every number in the pivot row by the pivot element, which makes the pivot element that we currently work with 1. That is the number on the diagonal of the matrix.

Hence, the pivot number in the first row is  $a_{11}$ . Moreover, we divide all of the elements in the pivot row by  $a_{11}$ , which makes the first pivot number equal to 1:

$$\text{Row}_1 = \frac{\text{Row}_1}{a_{11}}$$

Then the augmented matrix would look like this:

$$\left( \begin{array}{cc|c} 1 & \frac{a_{12}}{a_{11}} & \frac{b_1}{a_{11}} \\ a_{21} & a_{22} & b_2 \end{array} \right)$$

Doing this will make eliminating the elements below the pivot element easier.

## 1.4 Forward Elimination Continued

Now, we will make the number below our pivot number 0,  $a_{21}$ , and we will do this to all of the elements if there were more. This is done by the following:

$$\text{Row}_2 = \text{Row}_2 - a_{21} \times \text{Row}_1$$

Then the augmented matrix would look like:

$$\left( \begin{array}{cc|c} 1 & \frac{a_{12}}{a_{11}} & \frac{b_1}{a_{11}} \\ 0 & a'_{22} & b'_2 \end{array} \right)$$

with  $a'_{22}$  and  $b'_2$  being the changed numbers after the elimination

## 1.5 Backward Substitution

After we do the forward elimination, we will get the augmented matrix to be in upper triangular form. This makes it very simple to solve for the unknowns, as we start from the last row and then substitute the known variables into the equations we do not know. We keep doing this until we find all of the unknowns.

Form the example that we have been using to find the value of  $x_2$  we perform,

$$x_2 = \frac{b'_2}{a'_{22}}$$

Then, we substitute  $x_2$  into the equation of the first row of the augmented matrix, and this will let us find what  $x_1$  is, and thus, we will solve the system. Now, let us go through how the program would run.

$$x_1 = \frac{b_1 - a_{12}x_2}{a_{11}}$$

## 1.6 Programming Running Example

Now, we will provide an example of running the program and explain the output and the process of obtaining the solution.

### 1.6.1 Input

We first run `make` and then, `./math_matrix 2 2 2 1 solve print` in the command line, and then we get a random matrix  $A$  and vector  $B$ . Please note that for this example, we are looking at the main sections of the code and not the parts that handle errors to ensure that it is easy to understand:

$$A = \begin{pmatrix} 3.302049 & -2.470746 \\ -5.825511 & -9.356506 \end{pmatrix}, \quad B = \begin{pmatrix} 5.202234 \\ -6.052324 \end{pmatrix}$$

### 1.6.2 Initial Augmented Matrix

The augmented matrix that we get from appending  $B$  to  $A$ :

$$\text{Augmented Matrix} = \left( \begin{array}{cc|c} 3.302049 & -2.470746 & 5.202234 \\ -5.825511 & -9.356506 & -6.052324 \end{array} \right)$$

The code that dose this is below as we simply add the values to  $A$  and then append  $B$  at the end,

```
for (int row = 0; row < N1; row++) {
    for (int col = 0; col < M1; col++) {
        augmentedMatrix[row][col] = A[row][col];
    }
    augmentedMatrix[row][M1] = B[row][0];
}
```

### 1.6.3 Normalizing Row 1

Then, the first row is normalized by dividing each element in the pivot row by the pivot number, which is 3.302049:

$$\left( \begin{array}{cc|c} 1.000000 & -0.748246 & 1.575457 \\ -5.825511 & -9.356506 & -6.052324 \end{array} \right)$$

The code that does this is below. We iterate through the pivot row and divide every element by this first pivot value.

```
double pivot = augmentedMatrix[pivotRow][pivotRow];
for (int col = pivotRow; col <= M1; col++) {
    augmentedMatrix[pivotRow][col] /= pivot;
}
```

#### 1.6.4 Forward Elimination for Row 1

After normalization, the program begins the forward elimination to make the numbers below the pivot zero. The matrix then turns into the following:

$$\left( \begin{array}{cc|c} 1.000000 & -0.748246 & 1.575457 \\ 0.000000 & -13.715423 & 3.125515 \end{array} \right)$$

This step includes subtracting the correct multiple of the first row from the second row, and the code that does this is below:

```
for (int elimRow = pivotRow + 1; elimRow < N1; elimRow++) {
    double factor = augmentedMatrix[elimRow][pivotRow];
    for (int col = pivotRow; col <= M1; col++) {
        augmentedMatrix[elimRow][col] -= factor * augmentedMatrix[pivotRow][col];
    }
}
```

#### 1.6.5 Normalizing Row 2

Then the program normalizes the second row just like we did the first time, but now we are dividing the second pivot row by the pivot number, which is  $-13.715423$ , and we get:

$$\left( \begin{array}{cc|c} 1.000000 & -0.748246 & 1.575457 \\ 0.000000 & 1.000000 & -0.227883 \end{array} \right)$$

This ensures the pivot element of the second row is also 1.

#### 1.6.6 Back Substitution

Now, after the forward substitution is done, we have the augmented matrix in the upper triangular form:

$$\left( \begin{array}{cc|c} 1.000000 & -0.748246 & 1.575457 \\ 0.000000 & 1.000000 & -0.227883 \end{array} \right)$$

Now, the program starts the back substitution from the last row to solve for the variables. And the last row corresponds to the following equation  $0 \cdot x_1 + 1 \cdot x_2 = -0.227883$ , then we solve for  $x_2$ :

$$x_2 = -0.227883$$

Then the program moves to the first row which is  $1 \cdot x_1 - 0.748246 \cdot x_2 = 1.575457$ , and then the program substitutes the value of  $x_2$  to solve for  $x_1$ :

$$x_1 = 1.575457 + 0.748246 \times (-0.227883)$$

This results in:

$$x_1 = 1.404944$$

Then the program is done all of the steps, and the result is the solution to the system of linear equations, which is,  $Ax = B$  is:

$$x = \begin{pmatrix} 1.404944 \\ -0.227883 \end{pmatrix}$$

The code that performs the back substitution is below. It starts by initializing each  $x_i$  with the value from the augmented matrix corresponding to  $b_i$ . Then, it iteratively subtracts the known variables multiplied by their coefficients and then divides the coefficient of the unknown that is currently being solved if it is not normalized to 1 already.

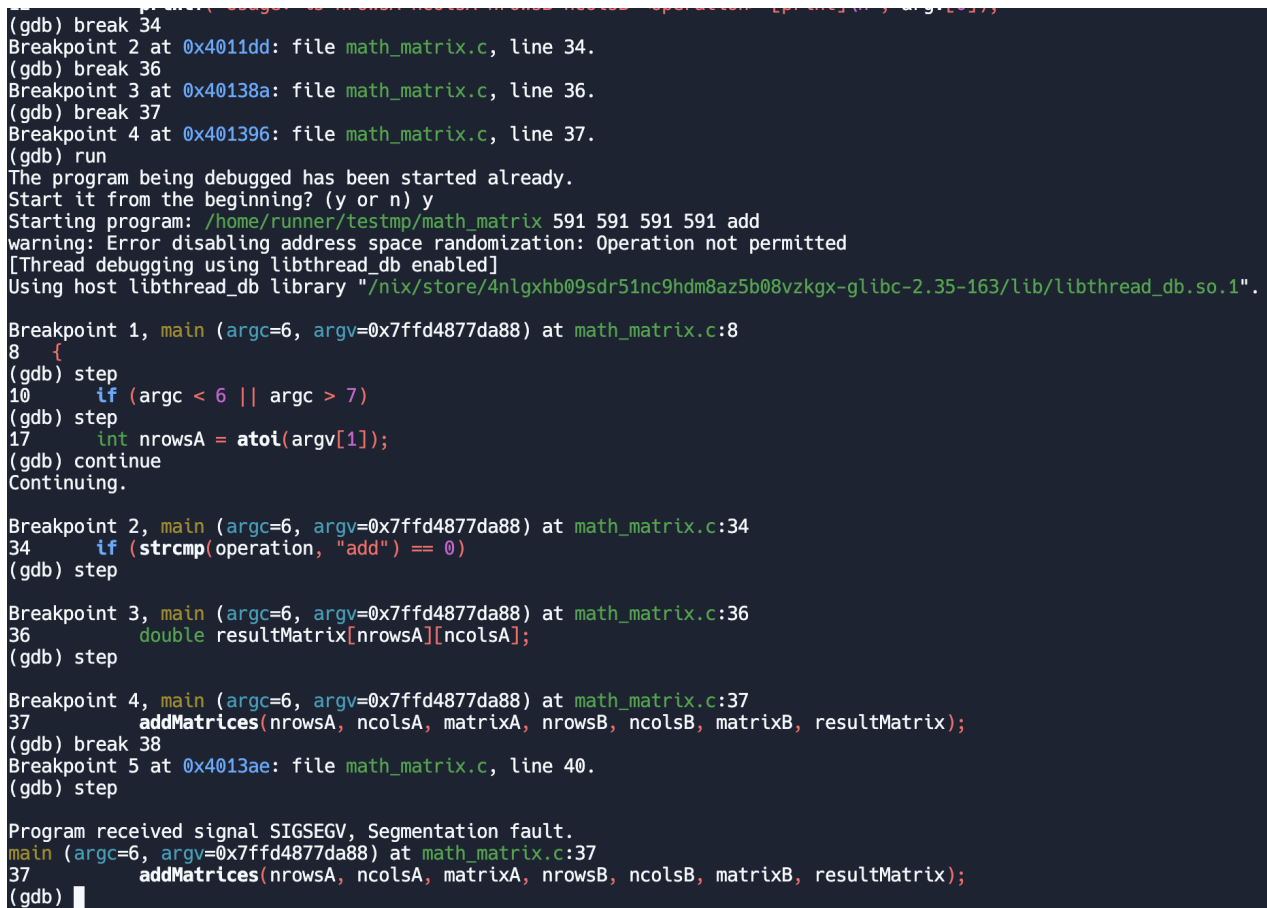
```

for (int row = N1 - 1; row >= 0; row--) {
    x[row] = augmentedMatrix[row][M1];
    for (int col = row + 1; col < N1; col++) {
        x[row] -= augmentedMatrix[row][col] * x[col];
    }
    if (augmentedMatrix[row][row] != 1) {
        x[row] /= augmentedMatrix[row][row];
    }
}
}

```

## 2 Segmentation Fault

Please note that I am on Mac OS using a MacBook M2 Pro, and using LLDB as GDB is not an option; due to this, debugging with LLDB is different than with GDB, and it needs to be more specific. I kept getting the following error, `queue = 'com.apple.main-thread', stop reason = EXC.BAD_ACCESS (code=2, address=0x16f603ff8)`. This was a common error for those trying to debug a segmentation fault across the internet using LLDB. Despite getting this error, I used Replit to try and use the GDB, and it worked, but it still needed to be specific. Despite this, I was able to get more information about the error. Now, let us explain what is happening.



```

(gdb) break 34
Breakpoint 2 at 0x4011dd: file math_matrix.c, line 34.
(gdb) break 36
Breakpoint 3 at 0x40138a: file math_matrix.c, line 36.
(gdb) break 37
Breakpoint 4 at 0x401396: file math_matrix.c, line 37.
(gdb) run
The program being debugged has been started already.
Start it from the beginning? (y or n) y
Starting program: /home/runner/testtmp/math_matrix 591 591 591 591 add
warning: Error disabling address space randomization: Operation not permitted
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/nix/store/4nlqxb09sdr51nc9hdm8az5b08vzkqx-glibc-2.35-163/lib/libthread_db.so.1".

Breakpoint 1, main (argc=6, argv=0x7ffd4877da88) at math_matrix.c:8
8  {
(gdb) step
10  if (argc < 6 || argc > 7)
(gdb) step
17  int nrowsA = atoi(argv[1]);
(gdb) continue
Continuing.

Breakpoint 2, main (argc=6, argv=0x7ffd4877da88) at math_matrix.c:34
34  if (strcmp(operation, "add") == 0)
(gdb) step

Breakpoint 3, main (argc=6, argv=0x7ffd4877da88) at math_matrix.c:36
36  double resultMatrix[nrowsA][ncolsA];
(gdb) step

Breakpoint 4, main (argc=6, argv=0x7ffd4877da88) at math_matrix.c:37
37  addMatrices(nrowsA, ncolsA, matrixA, nrowsB, ncolsB, matrixB, resultMatrix);
(gdb) break 38
Breakpoint 5 at 0x4013ae: file math_matrix.c, line 40.
(gdb) step

Program received signal SIGSEGV, Segmentation fault.
main (argc=6, argv=0x7ffd4877da88) at math_matrix.c:37
37  addMatrices(nrowsA, ncolsA, matrixA, nrowsB, ncolsB, matrixB, resultMatrix);
(gdb)

```

Figure 1: Segmentation fault, terminal GDB debugging

As we can see from the figure, the line causing the is 37, but since that is the Replit file, the line is off from the actual file, and in my VScode, the line is 41. If we were using GDB on Windows, we could get more information, but this tells us enough information, and we can break down why this is happening. We are calling `./math matrix 591 591 591 591 add`, which is creating a vast matrix that is 591\*591 sized matrix A and 591\*591 sized matrix B, and the program is trying to add these two matrices and stores them in a result matrix of the same size. The main issue is why we are getting the segmentation fault error issues with stack memory overflow. This is because of the allocation of a large matrix as local variables in the primary function exceed the stack size

limit of my environment. The allocated memory for the matrices in the `addMatrices` function is substantial. Specifically, for three matrices each of size  $591 \times 591$  doubles, the required memory is:

$$3 \times 591^2 \times 8 \text{ bytes}$$

With the notion that each double usually occupies 8 bytes, the total memory allocation is approximately 24 MB. This substantial allocation on the stack can easily lead to a segmentation fault due to exceeding the stack's memory limits. One potential solution is to use dynamic memory allocation, which allocates memory on the heap instead of the stack.

### 3 Running The Program

In order to run the program, follow the commands below:

1. Compile the program:

```
make
```

2. Run the program:

```
./math_matrix <nrowsA> <ncolsA> <nrowsB> <ncolsB> <operation> [print]
```

### 4 Appendix

---

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <string.h>
4
5  void generateMatrix(int rows, int cols, double matrix[rows][cols])
6  {
7      // Iterate over each row of the matrix
8      for (int i = 0; i < rows; i++)
9      {
10         // Within each row iterate over each column
11         for (int j = 0; j < cols; j++)
12         {
13             // Assign a random double from -10.0 to 10.0 to the current
14             // matrix element
15             matrix[i][j] = ((double)rand() / (double)(RAND_MAX)) * 20.0 -
16                             10.0;
17         }
18     }
19
20     // Prints a matrix with specified dimensions
21     void printMatrix(int rows, int cols, double matrix[rows][cols])
22     {
23         // Loop through each row.
24         for (int i = 0; i < rows; i++)
25         {
26             // Loop through each column in the current row
27             for (int j = 0; j < cols; j++)
28             {
29                 // Print the current element with 6 decimal places
30                 printf("%.6f ", matrix[i][j]);
31             }
32             printf("\n");
33     }
```

```

34
35 // Adds two matrices A and B and then stores result in the result matrix
36 void addMatrices(int N1, int M1, double A[N1][M1], int N2, int M2, double B[
    N2][M2], double result[N1][M1])
37 {
38     // Check if dimensions of both matrices are the same
39     if ((N1 == N2) && (M1 == M2))
40     {
41         for (int i = 0; i < N1; i++)
42         {
43             for (int j = 0; j < M1; j++)
44             {
45                 // Add corresponding elements and store in result
46                 result[i][j] = A[i][j] + B[i][j];
47             }
48         }
49     }
50     else
51     {
52         printf("Cannot add the matrices because of incompatible matrix sizes
            \n");
53     }
54 }
55
56 // Subtracts matrix B from matrix A and stores in a result matrix
57 void subtractMatrices(int N1, int M1, double A[N1][M1], int N2, int M2,
    double B[N2][M2], double result[N1][M1])
58 {
59     // check that matrices A and B have the same dimensions
60     if ((N1 == N2) && (M1 == M2))
61     {
62         for (int i = 0; i < N1; i++)
63         {
64             for (int j = 0; j < M1; j++)
65             {
66                 // Subtract element of B from A and store in result
67                 result[i][j] = A[i][j] - B[i][j];
68             }
69         }
70     }
71     else
72     {
73         printf("Cannot subtract the matrices due to incompatible sizes\n");
74     }
75 }
76
77 // Function to multiply two matrices and stores in result in result matrix
78 void multiplyMatrices(int N1, int M1, double A[N1][M1], int N2, int M2,
    double B[N2][M2], double result[N1][M2])
79 {
80     // checking to make sure that the matrix multiplication can be done
81     if (M1 == N2)
82     {
83         // iterate through the rows of the matrix
84         for (int rowA = 0; rowA < N1; rowA++)
85         {
86             // iterate through the cols of the matrix
87             for (int colB = 0; colB < M2; colB++)
88             {
89                 double sum = 0.0;
90                 /*

```

```

91         iterate along the rows and down the columns to perform the
           multiplication
92     by multiplying the values and then adding them up
93     */
94     for (int k = 0; k < M1; k++)
95     {
96         sum += A[rowA][k] * B[k][colB];
97     }
98     // assign the sum to the result matrix
99     result[rowA][colB] = sum;
100 }
101 }
102 }
103 else
104 {
105     printf("Cannot add the matrices because of incompatible matrix sizes
           \n");
106 }
107 }
108
109 void solveLinearSystem(int N1, int M1, double A[N1][M1], int N2, int M2,
           double B[N2][M2], double x[N1])
110 {
111     /*
112     Check if the dimesnsions are right to be able to sovle the linear system
           ,
113     if not then we will print a message saying that we cannot solve the
           system becuase of this
114     */
115     if (N1 == M1 && N1 == N2 && M2 == 1)
116     {
117         // initalize an augmented matrix that holds matrix A and the vector
           B
118         double augmentedMatrix[N1][M1 + 1];
119
120         // use a nested for loop to create the augmented matrix by adding
           the values to the matrix
121         for (int row = 0; row < N1; row++)
122         {
123             for (int col = 0; col < M1; col++)
124             {
125                 augmentedMatrix[row][col] = A[row][col];
126             }
127             // Setting b as the last column in the augmented matrix,
128             // after filling in the values from matrix A
129             augmentedMatrix[row][M1] = B[row][0];
130         }
131
132         // Forward elimination to put the augmented matrix in an upper
           triangular form
133         for (int pivotRow = 0; pivotRow < N1; pivotRow++)
134         {
135             // Check if the pivot element is 0 if it is 0 then there will be
           division by zero
136             // thus we cannot do that
137             if (augmentedMatrix[pivotRow][pivotRow] != 0)
138             {
139                 // Normalize the pivot row by dividing by the pivot element
140                 double pivot = augmentedMatrix[pivotRow][pivotRow];
141                 for (int col = pivotRow; col <= M1; col++)
142                 {

```



```

143         augmentedMatrix[pivotRow][col] /= pivot;
144     }
145
146     // Eliminate elements below the pivot to create zeros
147     for (int elimRow = pivotRow + 1; elimRow < N1; elimRow++)
148     {
149         double factor = augmentedMatrix[elimRow][pivotRow];
150         for (int col = pivotRow; col <= M1; col++)
151         {
152             augmentedMatrix[elimRow][col] -= factor *
                augmentedMatrix[pivotRow][col];
153         }
154     }
155 }
156 else
157 {
158     printf("Math error a zero pivot the solution may not be
        correct\n");
159 }
160 }
161
162 // Back substitution to solve for variables
163 for (int row = N1 - 1; row >= 0; row--)
164 {
165     // Initialize with the value from the augmented matrix
166     x[row] = augmentedMatrix[row][M1];
167     for (int col = row + 1; col < N1; col++)
168     {
169         // subtract the known variables
170         x[row] -= augmentedMatrix[row][col] * x[col];
171     }
172     // make sure that the row is normalized
173     if (augmentedMatrix[row][row] != 1)
174     {
175         x[row] /= augmentedMatrix[row][row];
176     }
177 }
178 }
179 else
180 {
181     printf("Incompatible dimensions.\n");
182 }
183 }

```

---

```

1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <string.h>
4  #include "functions.h"
5  #include <time.h>
6
7  int main(int argc, char *argv[])
8  {
9      // Check if the number of arguments is correct
10     if (argc < 6 || argc > 7)
11     {
12         printf("Usage: %s nrowA ncolA nrowB ncolB <operation> [print]\n"
            , argv[0]);
13         return 1;
14     }
15
16     // Creating variables based on the input from the command line

```

```

17  int nrowsA = atoi(argv[1]);
18  int ncolsA = atoi(argv[2]);
19  int nrowsB = atoi(argv[3]);
20  int ncolsB = atoi(argv[4]);
21  char *operation = argv[5];
22
23  // Seed the random number generator
24  srand(time(NULL));
25
26  // Initialize two sizes for matrix A and matrix B
27  double matrixA[nrowsA][ncolsA];
28  double matrixB[nrowsB][ncolsB];
29
30  // Fill in matrix A and matrix B with random values
31  generateMatrix(nrowsA, ncolsA, matrixA);
32  generateMatrix(nrowsB, ncolsB, matrixB);
33
34  // Checking which operation the user wants to perform (add)
35  if (strcmp(operation, "add") == 0)
36  {
37      // Initializing the size for the result matrix
38      double resultMatrix[nrowsA][ncolsA];
39
40      // Calling the add matrices function to add matrix A and B
41      addMatrices(nrowsA, ncolsA, matrixA, nrowsB, ncolsB, matrixB,
42                  resultMatrix);
43
44      /*
45      Check if the user wants to print the result matrix,
46      and that the matrix addition is possible, if so
47      then print out the matrix A and B and the result matrix
48      */
49      if (argc == 7 && strcmp(argv[6], "print") == 0 && (nrowsA == nrowsB)
50          && (ncolsA == ncolsB))
51      {
52          printf("Matrix A:\n");
53          printMatrix(nrowsA, ncolsA, matrixA);
54
55          printf("\nMatrix B:\n");
56          printMatrix(nrowsB, ncolsB, matrixB);
57
58          printf("\nResult of A + B:\n");
59          printMatrix(nrowsA, ncolsA, resultMatrix);
60      }
61  }
62  // Checking which operation the user wants to perform (subtract)
63  else if (strcmp(operation, "subtract") == 0)
64  {
65      // Initializing the size for the result matrix
66      double resultMatrix[nrowsA][ncolsA];
67
68      // Calling the function to subtract matrix A and B
69      subtractMatrices(nrowsA, ncolsA, matrixA, nrowsB, ncolsB, matrixB,
70                      resultMatrix);
71
72      /*
73      Check if the user wants to print the result matrix,
74      and that the matrix subtraction is possible, if so
75      then print out the matrix A and B and the result matrix
76      */
77      if (argc == 7 && strcmp(argv[6], "print") == 0 && (nrowsA == nrowsB)
78          && (ncolsA == ncolsB))

```

```

74     {
75         printf("Matrix A:\n");
76         printMatrix(nrowsA, ncolsA, matrixA);
77
78         printf("\nMatrix B:\n");
79         printMatrix(nrowsB, ncolsB, matrixB);
80
81         printf("\nResult of A - B:\n");
82         printMatrix(nrowsA, ncolsA, resultMatrix);
83     }
84 }
85 // Checking which operation the user wants to perform (multiply)
86 else if (strcmp(operation, "multiply") == 0)
87 {
88     // Initializing the side for the result matrix
89     double resultMatrix[nrowsA][ncolsB];
90
91     // call the multiply matrices function
92     multiplyMatrices(nrowsA, ncolsA, matrixA, nrowsB, ncolsB, matrixB,
93                     resultMatrix);
94
95     /*
96     Check if the user wants to print the result matrix,
97     and that the matrix addition is possible, if so
98     then print out the matrix A and B and the result matrix
99     */
100    if (argc == 7 && strcmp(argv[6], "print") == 0 && (ncolsA == nrowsB)
101        )
102    {
103        printf("Matrix A:\n");
104        printMatrix(nrowsA, ncolsA, matrixA);
105        printf("\nMatrix B:\n");
106        printMatrix(nrowsB, ncolsB, matrixB);
107
108        printf("\nResult of A * B:\n");
109        printMatrix(nrowsA, ncolsB, resultMatrix);
110    }
111 }
112 // Checking which operation the user wants to perform (solve)
113 else if (strcmp(operation, "solve") == 0)
114 {
115     //initialized vector to store the solution
116     double resultVector[nrowsA][1];
117
118     // Call the solveLinearSystem function
119     solveLinearSystem(nrowsA, ncolsA, matrixA, nrowsB, ncolsB, matrixB,
120                      resultVector);
121
122     /*
123     Check if the user wants to print the result and
124     that the input matrices are valid; if so, then print out the
125     matrices A and B and the result matrix
126     */
127     if ((argc == 7 && strcmp(argv[6], "print") == 0) && (nrowsA ==
128         ncolsA && nrowsA == nrowsB && ncolsB == 1))
129     {
130         printf("Matrix A:\n");
131         printMatrix(nrowsA, ncolsA, matrixA);
132
133         printf("\nMatrix B:\n");

```

```

130         printMatrix(nrowsB, ncolsB, matrixB);
131
132         // print out the solution
133         printf("\nResult of x=B/A:\n");
134         for (int i = 0; i < nrowsA; i++)
135         {
136             printf("%f\n", resultVector[i][0]);
137         }
138     }
139 }
140
141 else
142 {
143     printf("Invalid operation specified.\n");
144 }
145
146 return 0;
147 }

```

---

```

1  # Define the compiler
2  CC = gcc
3
4  # Define any compile-time flags
5  CFLAGS = -Wall -g
6
7  # Define the source file names
8  SRC = math_matrix.c functions.c
9
10 # Define the executable file name
11 EXEC = math_matrix
12
13 # Default target
14 all: $(EXEC)
15
16 $(EXEC): $(SRC)
17     $(CC) $(CFLAGS) -o $(EXEC) $(SRC)
18
19 # Clean target for removing compiled binary and object files
20 clean:
21     rm -f $(EXEC) *.o
22
23 # Phony targets for commands that do not represent files
24 .PHONY: all clean

```

---